

Model Comparison — ThePlatform.life AI (360° of Perspectives)

- **Best output quality** → **Claude Sonnet 4.6**. Noticeably more nuanced and emotionally honest — and it *shows* in the numbers: it writes ~80% more text per analysis than any other model. Cost ~6¢/analysis; it's also the slowest to *finish* (~76s) though the first card still appears in ~5s.
- **Best quality-for-speed** → **GPT-4.1**. Strong, accurate output at ~1.7¢/analysis and a ~19s full run — the best all-round paid-tier pick if Claude's latency is a concern.
- **Best cost** → **DeepSeek V4 Flash**. Solid, complete output for **under a tenth of a cent** per analysis. The obvious free-tier engine.
- **Keep Grok 4.3 as the default**. ~0.65¢/analysis, ~20s, dependable — a sensible middle ground.

Headline: every model produced all 9 cards correctly. Even the most expensive option is ~6¢ per analysis, and the cheapest is a rounding error.

Measured results (avg of 3 scenarios each)

Model	Provider	In / Out tokens	Time-to-first-card	Total latency	\$/analysis	\$/1,000	Output quality
Claude Sonnet 4.6	Anthropic	1,178 / 3,791	4.9s	76.1s	\$0.0604	\$60.40	★★★★★ most nuanced
GPT-4.1	OpenAI	1,137 / 1,998	5.3s	19.4s	\$0.0172	\$17.23	★★★★ accurate, slightly clinical
DeepSeek V4 Flash	DeepSeek	1,078 / 2,151	4.1s	24.3s	\$0.0007	\$0.67	★★★ capable, utilitarian
Grok 4.3 (current default)	xAI	1,178 / 2,074	7.3s	20.3s	\$0.0065	\$6.54	★★★ fast & creative

All five returned the full 9/9 cards on every run. Costs use each provider's list price applied to the measured token counts (table is illustrative — confirm rates against billing before contract).

What the numbers say

- **Token consumption is similar across models on the *input* side** (~1,080–1,180 — the same prompt) but **diverges sharply on output**. Claude generates ~3,800 output tokens of prose vs

~2,000–2,200 for the others — that single fact drives both its higher quality *and* its higher cost and latency.

- **Time-to-first-card (what the user actually feels) is ~4–7s for everyone.** Claude is among the *fastest* to show the first card (4.9s) even though it's slowest to finish — because our UI streams cards as they're ready, perceived speed stays good.
- **Total latency** favors GPT-4.1 (~19s) and Grok (~20s); DeepSeek ~24s; Claude ~76s (it's writing far more).

Cost at scale (per analysis × volume)

Model	99 analyses	1,000 analyses	10,000 analyses	100,000 analyses
Claude Sonnet 4.6	\$6	\$60	\$604	\$6,040
GPT-4.1	\$1.7	\$17	\$172	\$1,723
DeepSeek V4 Flash	\$0.067	\$0.67	\$6.70	\$67
Grok 4.3	\$0.65	\$6.5	\$65	\$654

Prompt caching (the static system prompt repeats every call) trims the input portion further after the first request — not modeled above, so real costs trend a little lower.

Recommendation

Model	Why
Claude Sonnet 4.6	Best-in-class nuance for a psychological product. ~6¢/analysis, ~5s to first card.
GPT-4.1	~90% of the quality at ¼ the latency and ¼ the cost — the pragmatic paid default.
DeepSeek V4 Flash	Complete, decent output for <\$0.001. Caps free-tier spend to near zero.
Grok 4.3	Keep — balanced cost/speed while we finalize the premium/free split.